

A Novel Heuristic for Directed Acyclic Graph Task Scheduling using Longest Betweenness Centrality

N.P. DAMINK, University of Twente, The Netherlands

Task scheduling is a well-known NP-hard problem that involves efficiently allocating computational tasks across available resources. Existing heuristic approaches, such as HEFT and MinMin, typically rely on local task properties and often fail to capture deeper structural dependencies between tasks. In this work, we propose a novel list-scheduling heuristic based on Longest Betweenness Centrality (LBC), a metric designed to quantify a task's influence by evaluating its presence on long dependency paths. We introduce several LBC-based ranking methods, including source-based and successor-weighted variants. The most effective variant, LBC-SRL, estimates task criticality by analyzing each task's current critical dependencies. Experimental evaluations across diverse synthetic DAGs demonstrate that LBC-SRL consistently outperforms classical heuristics such as MinMin, HCPT, and PEFT. Furthermore, real-world experiments validate these findings, with LBC-SRL even outperforming HEFT in the Epigenomics domain. These results highlight the value of incorporating global graph-theoretical metrics into task scheduling heuristics.

Additional Key Words and Phrases: Task Scheduling, Directed Acyclic Graphs (DAGs), Longest Betweenness Centrality (LBC), Parallel Computing, Node Degree, Makespan, Load Balancing

1 INTRODUCTION

Task scheduling is a fundamental problem in computer science with widespread applications in domains such as cloud computing [4], data processing [13] and real-time systems [19]. In the modern era, single processor systems are often insufficient for processing the huge amounts of data generated. This has led to the emergence of parallel and distributed computing environments, where multiple processors or computing nodes execute tasks concurrently [6]. Although these environments significantly improve computational power, they also introduce considerable complexity into the scheduling process. Task scheduling involves assigning computational tasks to available resources, such as multiple CPUs, distributed nodes, or virtual machines. The primary goal is optimizing performance metrics, such as minimizing makespan or balancing load distribution across processors. Suboptimal scheduling can result in resource under-utilization, increased execution times, and higher operational costs. In large-scale environments, even small inefficiencies can cause significant performance bottlenecks and financial waste. Therefore, effective task scheduling is of great importance in modern computing environments.

1.1 Background

In task scheduling, an application's components and their dependencies can be represented by a Directed Acyclic Graph (DAG). Formally,

a task DAG is a graph $G = (V, E)$, where $V = \{V_1, V_2, \dots, V_n\}$ is a finite set of tasks and E is a set of directed edges. We denote the number of tasks as $n = |V|$ and the number of edges as $e = |E|$. Each task $V_i \in V$ is associated with a processing cost c_i . This represents the time or computational effort required to execute the task. Each directed edge from task V_i to task V_j , denoted as $(V_i, V_j) \in E$, indicates that V_i must be completed before task V_j can begin. By definition, a DAG does not contain cycles, which means that there is no sequence of tasks $V_{i_1}, V_{i_2}, \dots, V_{i_k}, V_{i_1}$ such that each $(V_{i_l}, V_{i_{l+1}}) \in E$. Thus, the structure of a DAG directly represents the order in which tasks must be executed. DAG-based models are especially useful in parallel and distributed computing environments, where tasks can run concurrently if dependencies are respected.

One of the main objectives in task scheduling is minimizing the makespan. Makespan (C_{max}) is commonly used as a performance metric to evaluate scheduling algorithms. It indicates the completion time of the last task [20]. Effectively, it measures the total time required to complete a set of tasks in a scheduling problem. Note that this is different from the cumulative processing time of all tasks, as our parallel computing environment allows multiple tasks to execute concurrently. Formally, the makespan can be defined as

$$C_{max} = \max_{i \in \text{tasks}} C_i \quad (1)$$

Where C_i is the finishing time of task i and tasks is the set of all tasks to be scheduled.

We now introduce key definitions that will be used throughout the remainder of this paper.

Definition 1: Let $Pred(t_i)$ denote the set of direct predecessors of the task t_i . Formally, in the context of DAGs, this set is defined as:

$$Pred(t_i) = \{t_j \mid (t_j, t_i) \in E\} \quad (2)$$

Definition 2: Let $Succ(t_i)$ denote the set of direct successors of the task t_i . Formally, in the context of DAGs, this set is defined as:

$$Succ(t_i) = \{t_j \mid (t_i, t_j) \in E\} \quad (3)$$

We further define $Succ^*(t_i)$ as the set of all successors of t_i . That is, the set of all tasks $t_j \in V$ such that there exists a directed path from t_i to t_j in the DAG.

Definition 3: Let the *critical path* denote the longest path from an entry task to an exit task within the DAG. Formally, let the sets of entry and exit tasks be defined as:

$$T_{entry} = \{t_i \mid Pred(t_i) = \emptyset\} \quad (4)$$

$$T_{exit} = \{t_i \mid Succ(t_i) = \emptyset\} \quad (5)$$

Then, the *critical path* is the path of maximum total weight among all directed paths from any task in T_{entry} to any task in T_{exit} . This path determines the minimum completion time of the entire task DAG and forms the bottleneck for scheduling.

TSelT 43, July 4, 2025, Enschede, The Netherlands

© 2025 University of Twente, Faculty of Electrical Engineering, Mathematics and Computer Science.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

Definition 4: Let $EST(t_i, p_j)$ represent the earliest start time of task t_i on processor p_j :

$$EST(t_i, p_j) = \max \left\{ T_{available}(p_j), \max_{t_p \in Pred(t_i)} (AFT(t_p) + c_{p,i}) \right\} \quad (6)$$

where $T_{available}(p_j)$ is the earliest time at which processor p_j is ready for task execution, $AFT(t_p)$ is the actual finish time of predecessor t_p , and $c_{p,i}$ is the communication costs between tasks t_p and t_i . Note that $c_{p,i} = 0$ in our study, as communication costs are not considered in our model.

Definition 5: Let $EFT(t_i, p_j)$ represent the earliest finish time of task t_i on processor p_j :

$$EFT(t_i, p_j) = EST(t_i, p_j) + w_{i,j} \quad (7)$$

where $w_{i,j}$ is the processing cost of task t_i on processor p_j . However, we assume a homogeneous processor environment, meaning that each processor executes tasks with the same efficiency. Consequently, the cost of executing task t_i on any processor p_j is uniform and given by $w_{i,j} = c_i$.

1.2 Problem Statement

This study considers the task scheduling problem $P \mid prec \mid C_{max}$ in Graham's notation [14]. This notation represents a scheduling problem on parallel machines with identical processors (P). Each task is associated with a processing cost and is subject to precedence constraints ($prec$). Finally, the objective of this problem is to minimize the makespan (C_{max}). This study focuses specifically on static scheduling, where all task and dependency information is known in advance. Unlike dynamic scheduling, which makes decisions at runtime, static scheduling allows optimization based on the complete task DAG structure.

Since this problem is NP-complete [24], it is believed that it is not possible to create an algorithm that finds the optimal solution in polynomial time. Consequently, heuristic algorithms are essential for finding efficient and scalable solutions.

This paper investigates how the Longest Betweenness Centrality (LBC) metric, introduced in section 3.1, can effectively be integrated into heuristic task scheduling algorithms for parallel and distributed computing systems. To address this main research objective, we will investigate the following key questions:

- How do different LBC-based scoring strategies impact the makespan of task scheduling across diverse DAG structures?
- How does the computational complexity of the proposed LBC-based scheduling algorithms compare to that of state-of-the-art methods such as HEFT, PEFT, MinMin, and HCPT?
- How do the proposed LBC-based heuristics compare to state-of-the-art scheduling algorithms in terms of makespan performance across various task graph types and processor configurations?

To support the investigation of the proposed approach, the remainder of this paper is structured as follows. Section 2 reviews the state-of-the-art literature on heuristic scheduling algorithms and identifies key limitations in existing methods. Section 3 introduces the Longest Betweenness Centrality (LBC) metric and several variants of the novel LBC-based heuristic. Theoretical formulations and

complexity analyses are provided to establish their computational feasibility. Section 4 outlines the experimental methodology, including details on DAG generation and evaluation parameters. Section 5 presents the empirical results, comparing the proposed heuristics against established algorithms across a range of scenarios. Section 6 evaluates the effectiveness of the proposed heuristics on real-world DAGs obtained from diverse application domains. Finally, Section 7 discusses the findings, limitations, and directions for future research.

2 STATE-OF-THE-ART

A wide range of heuristic task scheduling algorithms have already been developed. State-of-the-art heuristics can be classified into three main categories: clustering, duplication-based, and list scheduling algorithms. Clustering algorithms aim to minimize communication costs by grouping related tasks. Duplication-based methods aim to reduce latency by replicating predecessors of tasks across processors. Finally, list scheduling algorithms assign priorities to tasks and schedule them in order, selecting the next task based on its priority. Among these classes, list scheduling algorithms have shown the best performance within the context of this study and will therefore be the primary focus of this research [2].

Well-known list-based scheduling algorithms are MinMin, HEFT, HCPT, and PEFT. These algorithms mainly rely on simple task graph properties, such as task execution time, immediate dependencies, or placement on the critical path. Although such properties can sometimes be effective for efficiently computing a near-optimal solution, they fail to exploit deeper structural characteristics of DAGs. As a result, these methods may overlook more optimal scheduling opportunities that can arise from global graph patterns. Therefore, a gap exists in leveraging advanced graph-theoretic concepts to decide scheduling decisions. This work aims to address this gap by proposing a novel heuristic that incorporates both global and local structural properties of DAGs to improve scheduling performance. The following section will provide a more detailed description of existing algorithms, which will serve as benchmarks in the comparative evaluation of our proposed approach.

2.1 MinMin

MinMin [15] is a greedy task scheduling heuristic that aims to minimize makespan by greedily assigning tasks that have the minimum EFT among all ready tasks. Specifically, for each ready task, it computes the EFT on all available processors and selects the task-processor pair that results in the minimum EFT. This process repeats until all tasks have been assigned.

2.2 HCPT

HCPT (Heterogeneous Critical Path Task scheduling) [16] consists of two phases, a ranking phase and a processor selection phase. In the ranking phase, tasks are prioritized based on their depth and position on the critical path. It schedules tasks in decreasing order of depth, giving precedence to those that are deep ancestors of critical-path nodes. In the processor selection phase, HCPT applies an EFT scheduling strategy. For each task, it selects the processor time slot that yields the earliest finish time while respecting precedence and availability constraints.

2.3 HEFT

HEFT (Heterogeneous Earliest Finish Time) [23] also follows these two main phases. In the ranking phase, tasks are prioritized based on their upward rank, which estimates the longest path from the task to the exit node. This metric can be interpreted as a reverse depth, highlighting tasks that lie on or near the critical path. During the processor selection phase, HEFT applies an EFT scheduling strategy similar to HCPT. However, HEFT also supports task insertion. This allows tasks to be scheduled in between existing time slots, provided that precedence and availability constraints are maintained.

2.4 PEFT

PEFT (Predict Earliest Finish Time) [3] also consists of two primary phases. In the ranking phase, tasks are prioritized using an Optimistic Cost Table (OCT). Each entry optimistically estimates the EFT by considering both computation and communication costs along critical paths. Tasks are then scheduled in descending order of their rank, where the rank is defined as the average value of task's entries in the OCT. In the processor selection phase, PEFT assigns tasks to processors by estimating the EFT based on predicted values. It uses a similar EFT insertion technique as HEFT. However, it improves task scheduling by using more accurate finish time predictions. These predictions account for communication delays more effectively, leading to better load balancing and increased efficiency.

3 PROPOSED METHOD: LBC-BASED SCHEDULING HEURISTICS

Algorithm 1 presents the structure of the proposed LBC-based scheduling heuristic, which is divided into two main phases. In the ranking phase, a supernode s^* is added using `ADD-SUPER-NODE(G)`, after which task scores are computed using one of the two LBC variants. These scores are then refined using a selected strategy from Section 3.2.1, applied through `APPLY-SCORING-STRATEGY($scores, G$)`. The introduced supernode is subsequently removed from the ranking before scheduling. In the processor selection phase, tasks are scheduled onto processors using an insertion-based EFT policy, implemented as `SCHEDULE($G, ranking, P, C$)`. The following sections provide a detailed explanation of the underlying theory behind the LBC prioritization metric, followed by a discussion of each phase in more detail.

Algorithm 1 Main pseudocode for the LBC-based heuristics.

```

function LBC-HEURISTIC(Graph  $G$ , Processors  $P$ , Costs  $C$ )
   $G \leftarrow \text{ADD-SUPER-NODE}(G)$ 
   $scores \leftarrow LBC_{(s)}(G, C)$ 
   $scores \leftarrow \text{APPLY-SCORING-STRATEGY}(scores, G)$ 
   $ranking \leftarrow \text{SCORE-GUIDED-RANKING}(G, scores)$ 
  remove  $s^*$  from  $ranking$ 
   $schedule \leftarrow \text{SCHEDULE}(G, ranking, P, C)$ 
  return  $schedule$ 
end function

```

3.1 Longest Betweenness Centrality

The scheduling heuristic is built around a novel metric called Longest Betweenness Centrality (LBC). Traditional betweenness centrality

measures the importance of a node within a network by counting how often it lies on the shortest paths between pairs of nodes [5]. In contrast, LBC redefines this concept by focusing on the longest paths. We define the longest path between two nodes t_i and t_j as the path with the maximum total weight among all directed paths from t_i to t_j . The existence of these paths is guaranteed by the acyclic nature of DAGs. The LBC for a node v can be formally defined as follows:

$$LBC(v) = \sum_{s \neq v \neq t \in V} \frac{L_{st}(v)}{L_{st}} \quad (8)$$

Where L_{st} is the number of longest paths from node s to node t and $L_{st}(v)$ is the number of those paths that pass through v .

Since makespan is directly affected by the length of the critical path, effective scheduling requires identifying tasks that frequently appear on it. The LBC metric captures this by prioritizing tasks that contribute the most to the DAG's critical structure. Unlike most state-of-the-art methods, LBC also considers the influence of tasks on long, non-critical paths, which can still significantly impact scheduling decisions. An illustrative example of how LBC is computed on a small weighted DAG is provided in Appendix C.

Brandes' algorithm [7] is widely recognized as one of the most efficient methods for computing standard betweenness centrality. We adapt a variation of this algorithm to compute LBC. Specifically, the new variant tracks maximum path lengths and counts the number of such longest paths that pass through each node. Moreover, the BFS or Dijkstra algorithm is replaced with a topological traversal to ensure that nodes are processed in the correct order for DAGs. This adaptation maintains the polynomial-time efficiency of the original algorithm, making it suitable for large task DAGs. The full pseudocode for this adaptation can be found in Appendix B.1.

In addition to the general LBC metric, we introduce a variant called Source-based Longest Betweenness Centrality (LBC_s). This version focuses on the influence of a node on the longest paths originating from entry tasks in the DAG. This approach still captures tasks that are critical to long execution chains, but avoids the computational overhead of evaluating all node pairs. As a result, the modification reduces the time complexity from $O(n(n+e))$ to $O(n+e)$. The LBC_s score for a node v is defined as:

$$LBC_s(v) = \sum_{v \neq t \in V, s \in T_{entry}} \frac{L_{st}(v)}{L_{st}} \quad (9)$$

The full pseudocode for this adaptation can be found in Appendix B.2.

3.2 Ranking Phase

3.2.1 Scoring Strategies for Task Prioritization. To determine the most effective use of LBC, this paper explores six task scoring approaches based on the metric. For each approach, a supernode is introduced into the DAG. This artificial node, denoted as s^* , is connected to all the source nodes in the original DAG. Formally, the new augmented DAG is defined as:

$$G' = (V \cup \{s^*\}, E \cup \{(s^*, v) \mid v \in V, \text{Pred}(v) = \emptyset\}) \quad (10)$$

In this augmented DAG, the new vertex s^* has directed edges to all nodes $v \in V$ for which $\text{Pred}(v) = \emptyset$. The main advantage of introducing this supernode is that the source nodes are included

in the longest path computations, allowing them to be properly quantified.

LBC: The first scoring approach directly uses the LBC score as the priority score for each node. The idea is that tasks with high LBC scores lie on many long dependency chains, contributing the most to the total execution time. Using this approach, the score of a node can be defined as:

$$Score(i) = LBC(i) \quad (11)$$

LBC Source-based (LBC-S): This second scoring approach uses the source-based variant of the LBC algorithm. Priority scores are computed analogously to the previous method, but using the source-based LBC metric:

$$Score(i) = LBC_s(i) \quad (12)$$

This approach only considers the longest paths that originate from the starting node, which is sufficient for our goal of minimizing these paths.

LBC Source-based with Direct Successors (LBC-SDS): This scoring approach builds upon the source-based variant of the LBC algorithm. However, the score also incorporates the scores of its dependent tasks. Specifically, the score is defined as a combination of the LBC_s score of the task and its immediate successors:

$$Score(i) = LBC_s(i) + \sum_{s \in Succ(i)} LBC_s(s) \quad (13)$$

The underlying rationale is to prioritize tasks that lead to successors with higher scores. This promotes the execution of tasks that contribute more significantly to the overall execution time.

LBC Source-based with All Successors (LBC-SAS): This scoring approach extends the previous algorithm by considering not only the immediate successors of a task, but all its transitive successors. Specifically, the score is computed as the sum of the LBC_s scores of the task itself and of all of its downstream tasks:

$$Score(i) = LBC_s(i) + \sum_{s \in Succ^*(i)} LBC_s(s) \quad (14)$$

This approach places even more emphasis on tasks that contribute to a larger number of high-scoring successors.

LBC Source-based with Weighted Successors (LBC-SWS): This approach further extends the previous algorithm by introducing a weighted scoring mechanism for the successor tasks. It draws inspiration from the work of Lin et al. [18], who proposed a weighted out-degree (WOD) metric to determine task priorities. This approach prioritizes tasks with a high out-degree and gives additional weight to those whose successor nodes have a low in-degree. Building on this concept, our approach adopts a similar, but inverted, weighting strategy. We will be prioritizing tasks that have a high in-degree, promoting early execution of tasks that enable the progression of more dependent tasks. The score is defined as follows:

$$Score(i) = LBC_s(i) + \sum_{s \in Succ^*(i)} LBC_s(s) \cdot |Prec(s)| \quad (15)$$

This approach emphasizes tasks that unlock access to highly connected successors, which are likely to become bottlenecks in later execution stages.

LBC Source-based Repeated Loop (LBC-SRL): This final scoring approach introduces an iterative variant of the source-based LBC

method. Instead of computing all scores in a single pass, LBC-SRL recomputes the LBC_s scores at each step of the ranking process. At each iteration, the algorithm identifies all source nodes and computes their source-based LBC scores using the current state of the DAG. The node with the highest score is selected, appended to the ranking, and then removed from the graph. This process is repeated until all nodes are scheduled. As this procedure involves iterative graph updates and score recalculations, it cannot be expressed as a simple formula. Instead, the complete pseudocode is provided in Appendix B.3. Although this design increases the computational complexity back to $O(n \cdot (n + e))$, it enables more informed scheduling decisions. At each step, the algorithm selects the task upon which the largest number of remaining tasks depend, thereby prioritizing the tasks that are most critical for subsequent execution stages.

In summary, these LBC-based strategies differ primarily in how they incorporate downstream influence: LBC and LBC-S focus solely on the centrality of the task, while the LBC-SDS, LBC-SAS, and LBC-SWS variants progressively expand the scope of influence from direct to weighted transitive successors. LBC-SRL stands out by recalculating the influence after each scheduling step, enabling prioritization that adapts to the evolving DAG structure.

3.2.2 Score-Guided Topological Scheduling. After computing the scores using one of the described scoring approaches, the final ranking is determined using a score-guided topological sort. During the traversal, tasks are ordered in a ranking that respects all precedence constraints. At each step, the algorithm selects the highest-scoring ready task using a priority queue implemented as a binary heap. Note that this ranking procedure is applied to all approaches except LBC-SRL, which already constructs the ranking iteratively during the score computation. The detailed pseudocode for this ranking procedure can be found in Appendix B.4.

3.3 Processor Selection Phase

The processor selection phase follows the task order of the generated ranking. This phase applies a similar insertion-based scheduling strategy used in HEFT [23]. Each task is scheduled on the processor that provides the lowest EFT. Task insertion between existing tasks is allowed, provided that dependency and availability constraints are satisfied. The detailed pseudocode for this scheduling procedure is provided in Appendix B.5.

3.4 Complexity Analysis

Table 1 presents the asymptotic time complexity of the scheduling algorithms evaluated in this study. In this table, we denote the number of tasks as $n = |V|$, the number of edges as $e = |E|$, and the number of processors as $p = |P|$. These complexity bounds are essential for evaluating the practical scalability of each approach, especially when applied to large task graphs.

The LBC and LBC-SRL heuristics have a worst-case time complexity of $O(n^2 \cdot p + n \cdot e)$. The term $O(n \cdot e)$ originates from the calculation of the LBC scores, with each LBC_s computation requiring $O(n + e)$ time. When repeated across n tasks, this results in a total of $O(n^2 + n \cdot e)$ time. However, the term $O(n^2)$ is asymptotically dominated by the processor selection phase, which has a

time complexity of $O(n^2 \cdot p)$. In this phase, each task is compared against all tasks on every processor to identify the optimal insertion point. Consequently, the total complexity for these two algorithms is $O(n^2 \cdot p + n \cdot e)$. In dense graphs, where $e \approx n^2$, the traversal term dominates, leading to a complexity of $O(n^3)$. In contrast, for sparser graphs or systems with many processors, the insertion phase dominates the overall complexity.

The remaining LBC variants have an overall complexity of $O(n^2 \cdot p)$. These approaches simplify the scheduling phase by using a lighter LBC_s with complexity $O(n + e)$ and a sorting step with $O(n \log n)$. Neither outweighs the cost of the insertion phase, which remains the asymptotically dominant component.

Similarly, the complexities of HEFT and PEFT are dominated by the insertion phase, resulting in a total complexity of $O(n^2 \cdot p)$.

For each scheduled task, MinMin evaluates all ready tasks across all processors to find the minimum EFT, leading to a time complexity of $O(n^2 \cdot p)$.

Finally, HCPT uses a simpler insertion-free selection mechanism, reducing its complexity to $O(n \cdot p)$. The dominant step for HCPT is the task ranking, which takes $O(n \log n)$ time.

Algorithms	Total Complexity
LBC, LBC-SRL	$O(n^2 \cdot p + n \cdot e)$
LBC-S, LBC-SDS, LBC-SAS, LBC-SWS, HEFT, MinMin, PEFT	$O(n^2 \cdot p)$
HCPT	$O(n \log n)$

Table 1. Total time complexity of scheduling algorithms.

4 EXPERIMENTAL SETUP

4.1 System Configuration

All experiments were carried out in a Windows 11 environment using the Windows Subsystem for Linux 2 (WSL 2) running Ubuntu 22.04.5 LTS with the Linux kernel version 5.15. The host machine was a Lenovo ThinkPad with an 11th-Gen Intel Core i7-11800H processor and 32GB of RAM. Python 3.10.12 was used for implementation and R 4.1.2 was used for statistical analysis and visualization. The baseline algorithms HEFT, HCPT and MinMin, used for the comparative evaluation, were adopted directly from the implementations by Canon et al. [8]. Each experimental configuration was executed 1000 times using randomly generated DAGs to account for variability. Details of the evaluation dataset, experimental parameters, and DAG generation techniques are described in the following section. All source code and additional resources are publicly available through the accompanying GitHub repository [10].

4.2 Evaluation Data

To evaluate the proposed algorithms, we use synthetically generated DAGs with $n = 100$ vertices. The processing costs are sampled uniformly over a range of 1 to 20. Although actual processing costs vary significantly by domain, this range provides a balanced set of task execution times, suitable for a general evaluation. The methods used for DAG generation produce DAGs with varying structural properties, allowing us to assess the strengths and limitations of

each heuristic across various DAG topologies. These generation techniques were implemented and analyzed in detail by Canon et al. [8]. The following methods will be used for evaluation:

Erdős-Rényi-Based Random Generation: This generation method is based on the Erdős-Rényi algorithm [11], which constructs DAGs by adding edges between pairs of vertices with independent probability p . To maintain acyclicity, edges are only added in the upper triangle of the adjacency matrix. The value of p controls the graph's sparsity or density. In our study, we use $p = 0.1$ and $p = 0.5$. The former produces sparse graphs with few edges between the nodes, while the latter produces dense graphs with many edges between the nodes. These p -values allow us to evaluate whether our algorithm performs better on sparse, dense, or both types of graphs.

Uniform Random DAG Generation: This generation method uniformly samples from the set of all labeled DAGs with n vertices using a recursive counting approach [21]. It guarantees an unbiased distribution within this well-defined class, ensuring every labeled DAG has the same probability of being selected. This enables an evaluation of the scheduling algorithms without any bias towards particular DAG shapes or structural patterns. However, the total number of dense DAGs significantly exceeds that of sparse ones in the sample space. This is due to the increasing number of possible edge combinations as graphs become denser. Consequently, dense graphs naturally appear more frequently in the generated dataset.

Random Orders Method: This method derives a DAG from randomly generated orders [25]. These orders are generated by intersecting k random total orders. The parameter k controls the depth and structure of the resulting DAGs. Higher values of k produce sparser and shallower DAGs, as edges are only added when multiple random total orders consistently agree on the direction. For our study, we select $k = 3$ to generate DAGs that reflect a balanced combination of task dependencies and parallelism. This choice produces graphs with moderate depth and sparsity, offering a realistic and challenging workload for scheduling algorithms.

Layer-by-Layer Method: The layer-by-layer method, first proposed by Adam et al. [1], is adapted in this study using a variant introduced by Canon et al. [8]. This method constructs DAGs by assigning vertices to layers and adding edges between layers with probability p . In this study, we use \sqrt{n} layers and set $p = 0.5$. This layered structure naturally models workflows with sequential phases or stages, which consequently enables a realistic assessment of scheduling algorithms under layered dependency constraints.

5 RESULTS

5.1 Comparative Evaluation of LBC-Based Ranking Strategies

To evaluate the effectiveness of the proposed LBC-based ranking strategies, the absolute makespan differences achieved by each approach across a diverse set of DAG topologies and processor configurations were measured. For each scheduling run, we compute the makespan C_{method} achieved by a given heuristic and subtract the minimal makespan C_{min} obtained across all heuristics for that specific configuration. Formally, the value plotted on the y-axis is:

$$\Delta C = C_{method} - C_{min} \quad (16)$$




Fig. 1. Comparative performance of LBC-based task ranking strategies across various DAG topologies and processor configurations. The x-axis shows the absolute makespan difference. Lower values reflect better scheduling performance. 'Erd-L' and 'Erd-B' refer to Erdős-Rényi DAGs with edge probabilities $p = 0.1$ and $p = 0.5$, respectively. Other methods include recursive, random order, and layered DAG generators.

Figure 1 summarizes these results. In this figure, algorithms positioned further to the left correspond to lower makespan values, indicating better scheduling performance. Among all ranking strategies, LBC-SRL consistently showed the best performance, demonstrating the effectiveness of recalculating task criticalities.

Furthermore, the figure shows that the graph density has a significant influence on the results. Sparser DAGs tend to produce a wider range of makespan values. These graphs have less restrictive precedence constraints, which offer greater flexibility in task ordering. In such cases, more informed heuristics can leverage the graph structure more effectively, resulting in lower makespan values. In contrast, as the processor count increases, denser DAGs fail to differentiate ranking strategies. These graphs restrict the range of available scheduling permutations more, causing the results to converge more to the length of the *critical path*.

Among the remaining ranking strategies, LBC-SWS achieves the second-best results, closely followed by LBC-SAS. These findings support the hypothesis that successor tasks play a significant role in determining the criticality of a task. This underscores the importance of incorporating successor information in the task prioritization process. In contrast, both LBC and LBC-S show consistently poor results across all DAG structures. This suggests that relying solely on a node's LBC score is insufficient for effective task prioritization.

Overall, LBC-SRL proved to be the most effective ranking strategy among those evaluated. Therefore, it is selected as the primary heuristic for the comparison against the established scheduling algorithms. However, because of its relatively lower computational complexity, LBC-SWS is also included in further analyses.

5.2 Evaluation Against Established Heuristics

Figure 2 presents the absolute difference in makespan between the established scheduling algorithms and the two proposed LBC-based heuristics. The figure clearly shows that HEFT consistently achieves the lowest makespan. This supports its reputation as one of the most effective list scheduling algorithms.

LBC-SRL achieves the second best results. It outperforms several well-known heuristics. For instance, it achieves significantly better results than both HCPT and MinMin. This outcome is unsurprising for MinMin, given its simplistic design, which does not consider the broader graph structure or future dependencies. HCPT also has a major limitation. Despite considering the graph structure, it does not utilize insertion-based scheduling. This reduces HCPT's flexibility in utilizing available processor time slots, resulting in less efficient schedules. This drawback is especially noticeable when the number of processors is low. Efficient task placement becomes more critical when the number of available processors is limited. Consequently, as shown in the figure, HCPT's performance improves with higher




Fig. 2. Absolute makespan differences for established scheduling heuristics and the proposed LBC-based algorithms (LBC-SRL, LBC-SWS) across various DAG generation models and processor counts. The x-axis shows the absolute makespan difference. Lower values reflect better scheduling performance.

processor counts, where its lack of insertion flexibility has less impact.

Furthermore, PEFT slightly underperforms compared to LBC-SRL. However, it should be noted that PEFT typically excels in environments with significant communication costs. Such costs are explicitly omitted in this study. This design decision likely reduces the advantage of PEFT’s optimistic cost table, which is optimized for both computation and communication costs. The advantage of LBC-SRL is especially visible in denser graphs (such as those generated using Erdős–Rényi model with $p = 0.5$ as well as DAGs generated using the recursive and layered methods). This is likely because LBC-SRL better captures global structural dependencies in complex DAGs, allowing it to prioritize tasks that are critical across long dependency chains.

LBC-SWS also performed well. It outperformed both HCPT and MinMin. This is likely, similar to LBC-SRL, due to its greater ability to account for global task dependencies and better task placement. It even performs similarly to PEFT. Although LBC-SWS does not quite match the performance of LBC-SRL, its lower computational complexity makes it a strong option in scenarios where efficiency is a priority.

Figure 3 provides a pairwise comparison of the heuristics, indicating how frequently each algorithm outperforms another. For each pair of heuristics h_i and h_j , we calculate the percentage of test cases where h_i produces a lower makespan than h_j . Formally, the win

percentage is defined as:

$$W(h_i, h_j) = \frac{\text{Number of cases where } C_{h_i} < C_{h_j}}{\text{Total number of comparisons}} \times 100 \quad (17)$$

where C_{h_i} and C_{h_j} are the makespan produced by heuristics h_i and h_j . This heatmap does not capture ties directly. However, these can be inferred by subtracting the total “win” percentages from 100% for each pair.

The pairwise results reveal that LBC-SRL beats PEFT in 33.2% of the cases, while being outperformed in only 21.7%. This clearly indicates a significant improvement. However, it only outperforms HEFT in 3.1% of cases and is outperformed in 42.0%, the remaining 54.9% being ties. This shows that, although LBC-SRL does not surpass HEFT, it remains competitive by tying more than half of the cases. Moreover, LBC-SRL is only outperformed in less than 2% of the cases by HCPT and MinMin. This demonstrates a significant improvement over these algorithms. Similarly, LBC-SWS performs well, beating PEFT in 13.9% of cases, tying 43.1% of the time and consistently outperforming both MinMin and HCPT.

6 REAL-WORLD COMPARISON

While synthetic DAGs allow controlled experimentation across a wide variety of structural configurations, they may not fully capture the complexity and irregularities present in real-world workloads. Consequently, to strengthen the practical relevance of our findings,




Fig. 3. Pairwise win percentage heatmap for all evaluated heuristics based on makespan performance. Each cell represents the proportion of runs in which the row algorithm outperforms the column algorithm.

we evaluated the proposed heuristics on real-world DAGs derived from three distinct application domains: biomedical sequencing, genomic analysis, and astronomical imaging. For each domain, we evaluated the heuristics on several DAGs. The exact DAG instances used for each domain are available in the supplementary GitHub repository [10]. Furthermore, all datasets are sourced from the Wflnstances collection maintained by WfCommons [9].

6.1 Epigenomics

The Epigenomics dataset consists of multiple DAGs derived from a genome analysis workflow [17]. Each DAG represents an execution instance that processes between 1 and 7 input files, with either 50,000 or 100,000 sequence entries per file. This variation results in DAGs containing 41 to 1,121 tasks, with execution costs ranging from 0.02 to 658.77 seconds. A detailed overview of the specific DAGs and their characteristics used in this domain’s evaluation of is provided in Appendix D.1.

At a high level, the DAGs are structured as parallel branches, one per input file. These branches converge at several points in the DAG, resulting in a combination of parallel and dependent regions.

Figure 4 presents the absolute makespan difference between heuristics across various processor counts in this domain. HCPT is excluded from this and subsequent comparisons due to consistently poor results, which would otherwise compress the plot scale and distort relevant differences.

LBC-SRL performs very well in this domain. It consistently achieves lower makespans than PEFT and MinMin. Notably, it even slightly outperforms HEFT in all tested situations. Although the median values are comparable, LBC-SRL shows narrower interquartile ranges, indicating more consistent and generally lower results. This suggests that its iterative criticality recalculation is particularly effective in workflows with repeated merging of parallel branches.

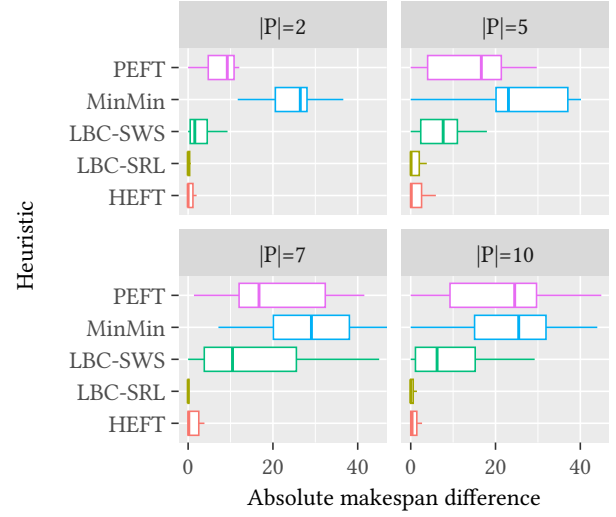


Fig. 4. Absolute makespan differences for established scheduling heuristics (PEFT, MinMin and HEFT) and the proposed LBC-based algorithms (LBC-SRL, LBC-SWS) on the Epigenomics workflow DAGs across varying processor counts. The x-axis shows the absolute makespan difference. Lower values indicate better scheduling performance.

LBC-SWS shows similar behavior to what was observed in the synthetic DAG evaluations, outperforming both PEFT and MinMin. This indicates that its successor-aware prioritization strategy generalizes well to real-world workflows.

6.2 Genome

The Genome dataset consists of multiple DAGs derived from a workflow based on the 1000 Genomes Project [12]. Each DAG corresponds to an execution instance that analyzes between 2 and 14 chromosomes, using either 100,000 or 250,000 sequences per chromosome. This variation results in DAGs containing 52 to 574 tasks, with execution costs ranging from 0.3 to 210 seconds. A detailed overview of the specific DAGs and their characteristics used in this domain’s evaluation of is provided in Appendix D.2.

Structurally, these DAGs have a repeated branching pattern, where each branch represents the analysis pipeline for a single chromosome. These branches are later merged, introducing both parallel and coordinated stages. A notable structural feature of this domain is the presence of mid-level entry tasks.

Figure 5 presents the absolute makespan difference between heuristics across various processor counts in this domain. Interestingly, the results reveal a notable shift in the performance of LBC-SRL. At lower processor counts, the results closely match those observed with synthetic DAGs. However, as the number of processors increases, the effectiveness of LBC-SRL declines. In contrast, the relative performance of the other heuristics remains largely consistent across processor counts.

This suggests that LBC-SRL may be less effective on graphs characterized by repeated branching patterns, particularly when more processors are available. This is likely due to the mid-level entry

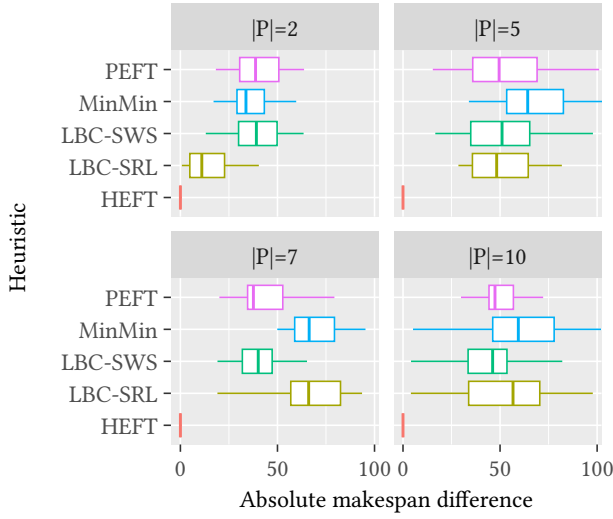


Fig. 5. Absolute makespan differences for established scheduling heuristics (PEFT, MinMin and HEFT) and the proposed LBC-based algorithms (LBC-SRL, LBC-SWS) on the Genome workflow DAGs across varying processor counts. The x-axis shows the absolute makespan difference. Lower values indicate better scheduling performance.

tasks. These lie on fewer longest paths from the supernode used in LBC-SRL’s ranking strategy. As a result, these tasks are assigned low priority and scheduled later than optimal, delaying downstream computation and increasing the makespan.

6.3 Montage

The Montage dataset consists of multiple DAGs derived from the Montage astronomical image mosaic workflow [22]. Each DAG corresponds to an execution instance that assembles images from the 2MASS survey to generate a mosaic of the sky. The instances vary by mosaic size, ranging from 0.05° to 0.5° in width and height. This results in DAGs containing between 58 and 1,738 tasks, with execution costs ranging from 0.02 to 45 seconds. A detailed overview of the specific DAGs and their characteristics used in this domain’s evaluation of is provided in Appendix D.3.

Structurally, the DAGs consist of three parallel branches. Within each branch, the number of tasks is determined by the target mosaic size. A final task merges the outputs from all three bands to produce a color composite, introducing a coordinated stage at the end of the workflow.

Figure 6 presents the absolute makespan difference between heuristics across various processor counts in this domain. These results closely align with the trends observed in the synthetic DAG evaluation. LBC-SRL outperforms all heuristics except HEFT, while LBC-SWS surpasses PEFT and MinMin. These findings suggest that the LBC-based heuristics generalize effectively to real-world workflows that combine parallel pipelines with a final synchronization stage.

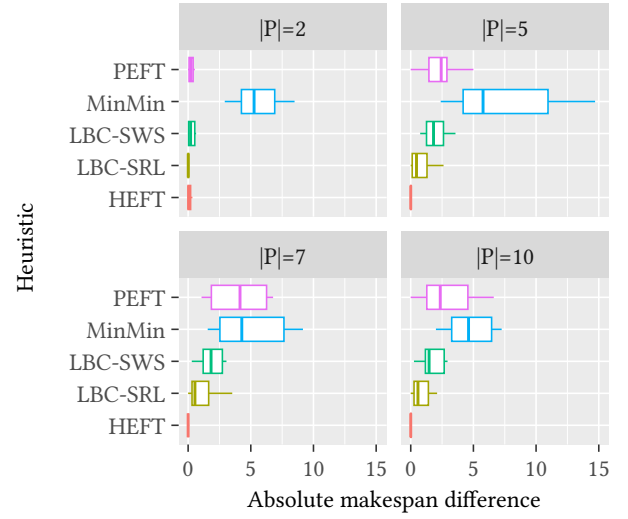


Fig. 6. Absolute makespan differences for established scheduling heuristics (PEFT, MinMin and HEFT) and the proposed LBC-based algorithms (LBC-SRL, LBC-SWS) on the Montage workflow DAGs across varying processor counts. The x-axis shows the absolute makespan difference. Lower values indicate better scheduling performance.

7 CONCLUSION

This paper introduced novel heuristics for task scheduling based on the Longest Betweenness Centrality (LBC) metric. We developed and evaluated several LBC-based ranking methods for task scheduling. Among them, the LBC-SRL variant demonstrated the most potential. Experimental results across a wide range of synthetic DAGs and processor configurations demonstrated that our approach yields competitive makespan performance compared to established heuristics, such as HEFT, MinMin, PEFT, and HCPT. The real-world evaluation confirmed the practical effectiveness of our approach, with particular strong performance in the Epigenomics workflow.

To address the first research question, we investigated how different LBC-based strategies influence scheduling performance. A comprehensive evaluation across multiple configurations was conducted. The results showed that LBC-SRL consistently achieved the lowest makespan, demonstrating the greatest potential compared to the other strategies.

Regarding computational efficiency, we observed that LBC-based heuristics have a time complexity of $O(n^2 \cdot p)$ or $O(n^2 \cdot p + n \cdot e)$. These complexities are comparable to established algorithms, such as HEFT and PEFT. Consequently, their performance remains practical for large task graphs, making it a practical alternative to existing heuristics without introducing significant computational overhead.

Finally, in comparison to established heuristics, LBC-SRL and LBC-SWS demonstrated competitive performance. Although HEFT remained the top performer in all synthetic DAG evaluations, LBC-SRL approached its results closely and outperformed MinMin, HCPT, and even PEFT in many configurations. Real-world evaluations

confirmed the validity of these results on practical DAGs. Notably, in the Epigenomics domain, LBC-SRL even outperformed HEFT. These findings highlight the effectiveness of LBC as a promising scheduling metric.

Overall, this research offers a novel contribution to the task scheduling domain by leveraging the global graph structure of task DAGs through the LBC metric. The proposed approach enables more informed scheduling decisions that reduce overall makespan. These findings indicate that scheduling strategies based on graph-theoretic centrality measures can effectively complement traditional heuristics in parallel and distributed computing environments. By integrating concepts from network analysis into task scheduling, this work opens a promising direction for future research on structure-aware heuristics that more fully exploit the topology of task graphs.

7.1 Limitations

Although the proposed LBC-based scheduling heuristics show promising results, several limitations should be acknowledged.

Assumption of Homogeneous Processing Resources: All experiments in the evaluation were conducted assuming homogeneous processors with identical capabilities. Although this simplifies the scheduling model, it may not reflect the heterogeneous nature of some real-world domains. However, the various heuristics can easily be extended to effectively handle heterogeneity. The scheduling phase of the proposed heuristics is based on the same EFT scheduling principle as HEFT, which is specifically designed to optimize heterogeneous environments.

Assumption of Static Scheduling Environment: This study assumes complete knowledge of the task graph and uses static scheduling. However, in many real-world systems, scheduling must be performed dynamically due to partial information, runtime variability, or unexpected delays. The current approach does not address these dynamic constraints.

Exclusion of Inter-Task Communication Costs: Communication delays between tasks assigned to different processors are excluded in our model. Communication overhead can significantly impact scheduling performance in distributed environments..

7.2 Future Work

Considering the findings and limitations of this study, there are several meaningful directions to explore for future research.

Integration of Communication Costs: Incorporating inter-task communication delays would provide a more realistic evaluation of scheduling performance, especially in distributed systems where data transfer between processors is non-negligible. Future adaptations of LBC-based heuristics may include adaptations that account for communication costs during the processor selection phase.

Extension to Heterogeneous Environments: Adapting the LBC framework to account for heterogeneous processors would increase its applicability to modern distributed computing platforms. This may involve modifying the processor selection strategy to account for processor-specific execution times per task. Moreover, the generation methods should be adapted to produce processing times that accurately reflect the characteristics of each processor.

Broader Algorithmic Benchmarking: While comparisons were made against several established heuristics, future work could expand this set to include more recent or domain-specific scheduling algorithms. This would provide a more comprehensive evaluation of where LBC excels and where it needs further improvement.

REFERENCES

- [1] Thomas L. Adam, K. M. Chandy, and J. R. Dickson. 1974. A comparison of list schedules for parallel processing systems. *Commun. ACM* 17, 12 (Dec. 1974), 685–690. <https://doi.org/10.1145/361604.361619>
- [2] Wakar Ahmad, Gaurav Gautam, Bashir Alam, and Bhoopesh Singh Bhati. 2024. An Analytical Review and Performance Measures of State-of-Art Scheduling Algorithms in Heterogeneous Computing Environment. *Archives of Computational Methods in Engineering* 31, 5 (2024), 3091–3113. <https://doi.org/10.1007/s11831-024-10069-8>
- [3] Hamid Arabnejad and Jorge G. Barbosa. 2014. List Scheduling Algorithm for Heterogeneous Systems by an Optimistic Cost Table. *IEEE Transactions on Parallel and Distributed Systems* 25, 3 (2014), 682–694. <https://doi.org/10.1109/TPDS.2013.57>
- [4] AR. Arunarani, D. Manjula, and Vijayan Sugumaran. 2019. Task scheduling techniques in cloud computing: A literature survey. *Future Generation Computer Systems* 91 (2019), 407–415. <https://doi.org/10.1016/j.future.2018.09.014>
- [5] David A. Bader, Shiva Kintali, Kamesh Madduri, and Milena Mihail. 2007. Approximating Betweenness Centrality. In *Algorithms and Models for the Web-Graph*, Anthony Bonato and Fan R. K. Chung (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 124–137.
- [6] D. Bertsekas and J. Tsitsiklis. 2015. *Parallel and Distributed Computation: Numerical Methods*. Athena Scientific, Athena. https://books.google.nl/books?id=n_Q5EAAAQBAJ
- [7] Ulrik Brandes. 2001. A faster algorithm for betweenness centrality*. *The Journal of Mathematical Sociology* 25, 2 (2001), 163–177. <https://doi.org/10.1080/0022250X.2001.9990249>
- [8] Louis-Claude Canon, Mohamad El Sayah, and Pierre-Cyrille Héam. 2019. A Comparison of Random Task Graph Generation Methods for Scheduling Problems. In *Euro-Par 2019: Parallel Processing*, Ramin Yahyapour (Ed.). Springer International Publishing, Cham, 61–73.
- [9] Tainã Coleman, Henri Casanova, Loïc Pottier, Manav Kaushik, Ewa Deelman, and Rafael Ferreira da Silva. 2022. WfCommons: A Framework for Enabling Scientific Workflow Research and Development. *Future Generation Computer Systems* 128 (2022), 16–27. <https://doi.org/10.1016/j.future.2021.09.043>
- [10] Niek Damink. 2025. LBC Task Scheduling Heuristic. https://github.com/Niek-Damink/LBC_TaskSched
- [11] P. Erdős and A. Rényi. 1959. On random graphs I. *Publ. math. debrecen* 6, 290–297 (1959), 18.
- [12] Rafael Ferreira da Silva, Rosa Filgueira, Ewa Deelman, Erola Pairo-Castineira, Ian M. Overton, and Malcolm P. Atkinson. 2019. Using simple PID-inspired controllers for online resilient resource management of distributed scientific workflows. *Future Generation Computer Systems* 95 (2019), 615–628. <https://doi.org/10.1016/j.future.2019.01.015>
- [13] Kannan Govindarajan, Supun Kamburugamuve, Pulasthi Wickramasinghe, Vibhatha Abeykoon, and Geoffrey Fox. 2017. Task Scheduling in Big Data - Review, Research Challenges, and Prospects. In *2017 Ninth International Conference on Advanced Computing (ICoAC)*. IEEE, Chennai, India, 165–173. <https://doi.org/10.1109/ICoAC.2017.8441494>
- [14] R.L. Graham, E.L. Lawler, J.K. Lenstra, and A.H.G. Rinnooy Kan. 1979. Optimization and Approximation in Deterministic Sequencing and Scheduling: a Survey. In *Discrete Optimization II*, P.L. Hammer, E.L. Johnson, and B.H. Korte (Eds.). Annals of Discrete Mathematics, Vol. 5. Elsevier, Cham, 287–326. [https://doi.org/10.1016/S0167-5060\(08\)70356-X](https://doi.org/10.1016/S0167-5060(08)70356-X)
- [15] Oscar H. Ibarra and Chul E. Kim. 1977. Heuristic Algorithms for Scheduling Independent Tasks on Nonidentical Processors. *J. ACM* 24, 2 (April 1977), 280–289. <https://doi.org/10.1145/322003.322011>
- [16] Jan Janeczek and Tarek Hagras. 2003. A Simple Scheduling Heuristic for Heterogeneous Computing Environments. In *Parallel and Distributed Computing, International Symposium on*. IEEE Computer Society, Los Alamitos, CA, USA, 104. <https://doi.org/10.1109/ISPDC.2003.1267650>
- [17] Gideon Juve, Ann Chervenak, Ewa Deelman, Shishir Bharathi, Gaurang Mehta, and Karan Vahi. 2013. Characterizing and profiling scientific workflows. *Future Generation Computer Systems* 29, 3 (2013), 682–692. <https://doi.org/10.1016/j.future.2012.08.015> Special Section: Recent Developments in High Performance Computing and Security.
- [18] Han Lin, Ming-Fan Li, Cheng-Fan Jia, Jun-Nan Liu, and Hong An. 2019. Degree-of-Node Task Scheduling of Fine-Grained Parallel Programs on Heterogeneous Systems. *Journal of Computer Science and Technology* 34, 5 (2019), 1096–1108. <https://doi.org/10.1007/s11390-019-1962-4>
- [19] C. L. Liu and James W. Layland. 1973. Scheduling Algorithms for Multiprogramming in a Hard-Real-Time Environment. *J. ACM* 20, 1 (Jan. 1973), 46–61. <https://doi.org/10.1145/321738.321743>
- [20] Michael L. Pinedo. 2016. *Scheduling: Theory, Algorithms, and Systems* (5 ed.). Springer Cham, Cham. <https://doi.org/10.1007/978-3-319-26580-3>
- [21] Robert W Robinson. 1973. Counting unlabeled acyclic digraphs. *New directions in the theory of graphs* (1973), 239–273.
- [22] M. Rynge, G. Juve, J. Kinney, J. Good, B. G. B., A. Merrihew, and E. Deelman. 2013. Producing an Infrared Multiwavelength Galactic Plane Atlas using Montage, Pegasus and Amazon Web Services. In *Proceedings of the 23rd Annual Astronomical Data Analysis Software and Systems (ADASS) Conference*. ADASS XXIII.
- [23] H. Topcuoglu, S. Hariri, and Min-You Wu. 2002. Performance-effective and low-complexity task scheduling for heterogeneous computing. *IEEE Transactions on Parallel and Distributed Systems* 13, 3 (2002), 260–274. <https://doi.org/10.1109/71.993206>
- [24] J.D. Ullman. 1975. NP-complete scheduling problems. *J. Comput. System Sci.* 10, 3 (1975), 384–393. [https://doi.org/10.1016/S0022-0000\(75\)80008-0](https://doi.org/10.1016/S0022-0000(75)80008-0)
- [25] Peter Winkler. 1985. Random orders. *Order* 1 (1985), 317–331.

APPENDICES

A AI STATEMENT

During the preparation of this work the author used ChatGPT in order to enhance the writing of this paper. After using this tool, the author reviewed and edited the content as needed and takes full responsibility for the content of the work.

B PSEUDOCODE

This appendix presents the detailed pseudocode for the algorithms described in the main text. The pseudocode is designed to clearly explain the algorithmic steps and key computations without relying on any specific programming language.

B.1 LBC Pseudocode

Algorithm 2 LBC Pseudocode.

```

function LBC(Graph G)
  LBC[v]  $\leftarrow$  0,  $v \in G.V$ ;
  topological_order  $\leftarrow$  TOPOLOGICAL-SORT(s);
  while topological_order  $\neq \emptyset$  do
    s  $\leftarrow$  topological_order[0]
    visited  $\leftarrow$  [s]
    P[w]  $\leftarrow$  [],  $w \in G.V$ ;
     $\sigma[t] \leftarrow 0$ ,  $t \in G.V$ ;  $\sigma[s] \leftarrow 1$ 
     $d[t] \leftarrow -1$ ,  $t \in G.V$ ;  $d[s] \leftarrow 0$ 
    for all v  $\in$  topological_order do
      if v  $\neq$  s and  $\sigma[v] = 0$  then
        skip to next iteration
      end if
      append v  $\rightarrow$  visited
      for all outgoing neighbors w of v do
        if  $d[w] < d[v] + c_v$  then
           $d[w] \leftarrow d[v] + c_v$ 
           $\sigma[w] \leftarrow \sigma[v]$ 
          P[w]  $\leftarrow$  [v]
        else if  $d[w] = d[v] + c_v$  then
           $\sigma[w] \leftarrow \sigma[w] + \sigma[v]$ 
          append v  $\rightarrow$  P[w]
        end if
      end for
    end for
     $\delta[t] \leftarrow 0$ ,  $t \in$  visited
    for all w  $\in$  REVERSED(visited) do
      for all v  $\in$  P(w) do
         $\delta[v] \leftarrow \delta[v] + \frac{\sigma[v]}{\sigma[w]} \cdot (1 + \delta[w])$ 
      end for
      if w  $\neq$  s then
        LBC[w]  $\leftarrow$  LBC[w] +  $\delta[w]$ 
      end if
    end for
    pop s  $\rightarrow$  topological_order
  end while
end function

```

B.2 Source-based LBC Pseudocode

The pseudocode is identical to the general LBC algorithm with one key addition: at the start of each iteration over the topological order, the algorithm checks if the current node has incoming neighbors. If it does, the node is skipped to ensure only source nodes initiate the traversal. Consequently, the following line needs to be added to the start of the main loop:

Algorithm 3 Early-exit condition to restrict traversal to source nodes only.

```

if incoming_neighbors of v  $\neq \emptyset$  then
  return LBC
end if

```

B.3 LBC-SRL Pseudocode:

Algorithm 4 Task Ranking using LBC Scores.

```

function LBC-SRL-RANKING(Graph G)
  ranking  $\leftarrow$  []
  while G.V  $\neq \emptyset$  do
    G  $\leftarrow$  ADD-SUPER-NODE(G)
    scores  $\leftarrow$  LBCs(G)
    G  $\leftarrow$  REMOVE-SUPER-NODE(G)
    ready_tasks  $\leftarrow$  {v  $\in$  G.V | in_degree(v) = 0}
    task  $\leftarrow$  maxv  $\in$  ready_tasks (scores[v])
    append task  $\rightarrow$  ranking
    REMOVE-NODE(G, task)
  end while
  return ranking
end function

```

B.4 Score-Guided Topological Sorting Pseudocode

Algorithm 5 Score-Guided Task Ranking Using LBC Scores.

```

function SCORE-GUIDED-RANKING(Graph G, Scores S)
  ranking  $\leftarrow$  []
  ready_tasks  $\leftarrow$  {v  $\in$  G.V | in_degree(v) = 0}
  heap  $\leftarrow$  construct heap for [S[v], v]  $\in$  ready_tasks
  while heap  $\neq \emptyset$  do
    v  $\leftarrow$  highest-priority node from heap
    append v to ranking
    for all successors w of v do
      remove edge (v, w) from G
      if w now has no incoming edges then
        insert w into heap with priority S[w]
      end if
    end for
  end while
  return ranking
end function

```

B.5 Insertion-based EFT Scheduling Pseudocode

Algorithm 6 Insertion-Based EFT Scheduling.

```

function SCHEDULE(Graph  $G$ , Ranking  $R$ , Processors  $P$ , Costs  $C$ )
   $task\_finish[t] \leftarrow 0, t \in G.V$ 
   $schedule \leftarrow []$ 
   $proc\_slots[p] \leftarrow [], p \in P$ 
  for all  $t \in R$  do
     $ready\_at \leftarrow \max\{task\_finish[d] \mid d \in predecessors(t)\}$ 
     $EFT \leftarrow \infty$ 
    for all  $p \in P$  do
       $start\_time \leftarrow \text{find first gap after } ready\_at \text{ on } p$ 
       $finish\_time \leftarrow start\_time + C[t]$ 
      if  $finish\_time < EFT$  then
         $best\_proc \leftarrow p$ 
         $EST \leftarrow start\_time$ 
         $EFT \leftarrow finish\_time$ 
      end if
    end for
    Reserve slot  $[start, end]$  on processor  $best\_proc$ 
     $task\_finish[t] \leftarrow EFT$ 
    Append  $\{task : t, proc : best\_proc, start : EST, end : EFT\} \rightarrow schedule$ 
  end for
  return  $schedule$ 
end function

```

C ILLUSTRATIVE EXAMPLE OF LONGEST BETWEENNESS CENTRALITY

This Appendix provides a concrete example demonstrating how the Longest Betweenness Centrality (LBC) metric is computed on a small weighted DAG. This example is intended for readers who need further clarification beyond the formal definition provided in the main text or for those who wish to develop a more intuitive understanding of the metric through a step-by-step application.

LBC Example:

The DAG, shown in Figure 7, contains six nodes and seven directed edges, each annotated with a weight that represents the execution time of the task.




Fig. 7. Example DAG with node colors representing LBC scores. Darker nodes indicate higher centrality, reflecting more frequent occurrence on longest weighted paths.

1: Longest Paths from Source A to Sink F

The LBC metric quantifies how often a node lies on the longest weighted paths between pairs of nodes in the DAG. Consequently, we must consider all longest paths between all source-destination pairs in the DAG. We will first consider the contribution to the LBC from the longest path between A and F

- (1) $A \rightarrow B \rightarrow D \rightarrow F$ with total weight: $2 + 4 + 4 = 10$
- (2) $A \rightarrow C \rightarrow D \rightarrow F$ with total weight: $2 + 4 + 4 = 10$
- (3) $A \rightarrow C \rightarrow E \rightarrow F$ with total weight: $2 + 4 + 5 = 11$

The longest path among these is the third, so only nodes C and E contribute to LBC for this path.

2: Additional Longest Paths From Node A To B, C, D and E

We now consider all longest paths originating from node A , assigning points to all nodes that lie along these paths. When multiple longest paths exist, the points are evenly distributed across all such paths.

- To D :
 - $A \rightarrow B \rightarrow D$ and $A \rightarrow C \rightarrow D$. Both are equal-length longest paths, so B and C each receive 0.5 points.
- To E :
 - $A \rightarrow C \rightarrow E$. Node C gains 1 point.
- To B or C : Paths to B and C are direct from A , so no intermediate nodes contribute to their LBC.

Thus, only considering all longest paths from source node A , the intermediate LBC scores are described in the following table:

Node	LBC Score
A	0
B	0.5
C	2.5
D	0
E	1
F	0

3: Computing the Final LBC Scores

To compute the final LBC scores, the same process is repeated for all valid source-target pairs in the DAG. The cumulative scores obtained are shown in the following table.

Node	LBC Score
A	0
B	0.5
C	2.5
D	1
E	2
F	0

This example illustrates how LBC identifies structurally significant nodes that frequently appear along critical long execution paths in the DAG.

D DATASET DETAILS

This appendix provides detailed summaries of the real-world DAGs used in the real-world evaluation, organized by dataset: Epigenomics, Genome, and Montage. Each table lists key statistics such as task counts, dependency counts, average degree, and cost metrics.

D.1 Epigenomics

Name	#Tasks	#Dependencies	Avg In/Out Degree	Max Cost	Min Cost	Avg Cost
epigenomics-chameleon-hep-1seq-100k-001	41	48	1.17	59.72	0.15	13.15
epigenomics-chameleon-hep-1seq-50k-001	73	88	1.21	61.07	0.10	17.04
epigenomics-chameleon-hep-2seq-100k-001	119	144	1.21	88.62	0.38	24.36
epigenomics-chameleon-hep-2seq-50k-001	223	274	1.23	70.00	0.08	16.29
epigenomics-chameleon-hep-3seq-100k-001	233	285	1.22	86.02	0.03	22.88
epigenomics-chameleon-hep-3seq-50k-001	445	550	1.24	76.50	0.03	18.09
epigenomics-chameleon-hep-4seq-100k-001	347	426	1.23	122.64	0.09	28.78
epigenomics-chameleon-hep-4seq-50k-001	671	831	1.24	84.15	0.07	20.62
epigenomics-chameleon-hep-5seq-100k-001	421	517	1.23	116.41	0.09	29.76
epigenomics-chameleon-hep-5seq-50k-001	817	1012	1.24	97.30	0.02	24.8
epigenomics-chameleon-hep-6seq-100k-001	507	623	1.23	481.98	0.07	26.07
epigenomics-chameleon-hep-6seq-50k-001	983	1218	1.24	73.11	0.06	18.36
epigenomics-chameleon-hep-7seq-100k-001	577	709	1.23	581.65	0.12	27.57
epigenomics-chameleon-hep-7seq-50k-001	1121	1389	1.24	658.77	0.03	20.12

D.2 Genome

Name	#Tasks	#Dependencies	Avg In/Out Degree	Min Cost	Max Cost	Avg Cost
1000genome-chameleon-10ch-100k-001	260	380	1.46	148.15	0.29	61.66
1000genome-chameleon-10ch-250k-001	410	530	1.29	141.39	1.40	63.36
1000genome-chameleon-12ch-100k-001	312	456	1.46	170.92	0.56	58.79
1000genome-chameleon-12ch-250k-001	492	636	1.29	151.85	1.65	57.18
1000genome-chameleon-14ch-100k-001	364	532	1.46	186.02	0.31	64.29
1000genome-chameleon-14ch-250k-001	574	742	1.29	156.88	0.47	61.61
1000genome-chameleon-2ch-100k-001	52	76	1.46	112.04	0.31	53.29
1000genome-chameleon-2ch-250k-001	82	106	1.29	105.00	1.93	54.10
1000genome-chameleon-4ch-100k-001	104	152	1.46	165.67	2.09	82.79
1000genome-chameleon-4ch-250k-001	164	212	1.29	165.63	0.38	72.47
1000genome-chameleon-6ch-100k-001	156	228	1.46	156.00	1.31	69.57
1000genome-chameleon-6ch-250k-001	246	318	1.29	182.06	0.85	70.47
1000genome-chameleon-8ch-100k-001	208	304	1.46	211.81	0.47	79.89
1000genome-chameleon-8ch-250k-001	328	424	1.29	186.58	0.35	66.22

D.3 Montage

Name	#Tasks	#Dependencies	Avg In/Out Degree	Min Cost	Max Cost	Avg Cost
montage-chameleon-2mass-005d-001	58	114	1.97	18.83	0.09	3.82
montage-chameleon-2mass-015d-001	310	798	2.57	18.66	0.05	2.76
montage-chameleon-2mass-01d-001	103	231	2.24	17.32	0.05	3.52
montage-chameleon-2mass-025d-001	619	1641	2.65	16.34	0.03	2.14
montage-chameleon-2mass-03d-001	748	1992	2.66	18.44	0.03	2.34
montage-chameleon-2mass-04d-001	1312	3540	2.70	18.32	0.02	2.30
montage-chameleon-2mass-05d-001	1738	4698	2.70	44.77	0.04	5.00